Product Owner case study · KYC onboarding

KYC approval fell 5.4 points in a month. Most of the arithmetic is known; most of the causes are not yet.

What the 2023 extracts show, what remains a hypothesis, and what to do first.

Data: KYC_summary and KYC_details extracts, 52,075 recorded decisions, Jul 25 – Sep 22, 2023 (prior period Jul 25 – Aug 23; recent period Aug 24 – Sep 22). Vendor research dated Oct 13, 2026. Case-study submission; it contains fraud-control and compliance findings from the extract.

The answer in seven points

1	What is measurable	Approval among recorded decisions, by day, document country and failing check. Not measurable: starts, completion, abandonment, retries, review outcomes, time to completion.
2	What changed	Approval 88.7% → 83.3% (−5.4 pp). Approvals −2,945: −1,587 from fewer decisions, −1,358 from the lower rate.
3	Arithmetic drivers	Fraud signals −2.05 pp, execution errors −1.62 pp, selfie / liveness −0.95 pp, no document type −0.60 pp. Face match improved (+0.36 pp).
4	Still hypotheses	1,057 failures (range 1,055–1,130) sit above baseline in four date windows. Fraud signals are not adjudicated; the execution errors are unattributed; vendor incidents on Sep 1 and 6 are documented, but impact on ARQ is not.
5	First (Phase 0, weeks 0–2)	Instrument the funnel; hold and retry execution errors; alert per reason code; reconcile control findings; adjudicate fraud samples; diagnose Mexican "no document type" with the vendor.
6	Needed	Event schema with privacy sign-off; vendor case access and incident reports; Risk and Compliance time; one medium and several small engineering items (sizes to confirm).
7	Verification	Acceptance criteria per action; experiments on eventual approval per person within 7 days. Scenarios for recurring causes: +0.9 pp to +4.3 pp, not forecasts. Confirmed control gaps could cost up to 2.3 pp.

We can measure approval among recorded decisions, not the start-to-finish funnel

Headline metric

Approval rate = approved (PASSED or APPROVED) ÷ recorded decisions.

Prior period 88.7%, recent 83.3%.

WARNING counts as not approved and not as a confirmed rejection. Its later review outcome is not in the data.

Each record is one user; the extract does not say whether it is a first, latest or final attempt, nor what event its date marks.

Metric	Status	Prior	Recent
Approval rate among recorded decisions	**Available**	88.7%	83.3%
Unresolved warning rate	**Available**	0.7%	1.1%
Completion rate (start to terminal state)	**Unavailable**	—	—
Eventual approval within 7 days	**Unavailable**	—	—
Abandonment rate	**Unavailable**	—	—
Retry rate and retry approval	**Unavailable**	—	—
Manual-review outcomes	**Unavailable**	—	—
Time to completion	**Unavailable**	—	—

To measure the funnel, log

- kyc_started
- document_selected
- document_captured
- selfie_captured
- kyc_submitted
- vendor_result
- kyc_decided
- review_completed
- account_enabled

with person, session and attempt ids, timestamps and app/SDK version.

Do not read the approval rate as completion: people who never submit, and approvals after retry or review, are invisible here. Whether true completion is lower or higher cannot be determined from these extracts.

Approval fell from 88.7% to 83.3%: 2,945 fewer approvals

88.7% → 83.3%

Approval among recorded decisions, prior → recent period (−5.4 pp)

−2,945

Approvals: −1,587 from fewer decisions at the prior rate, −1,358 from the lower rate

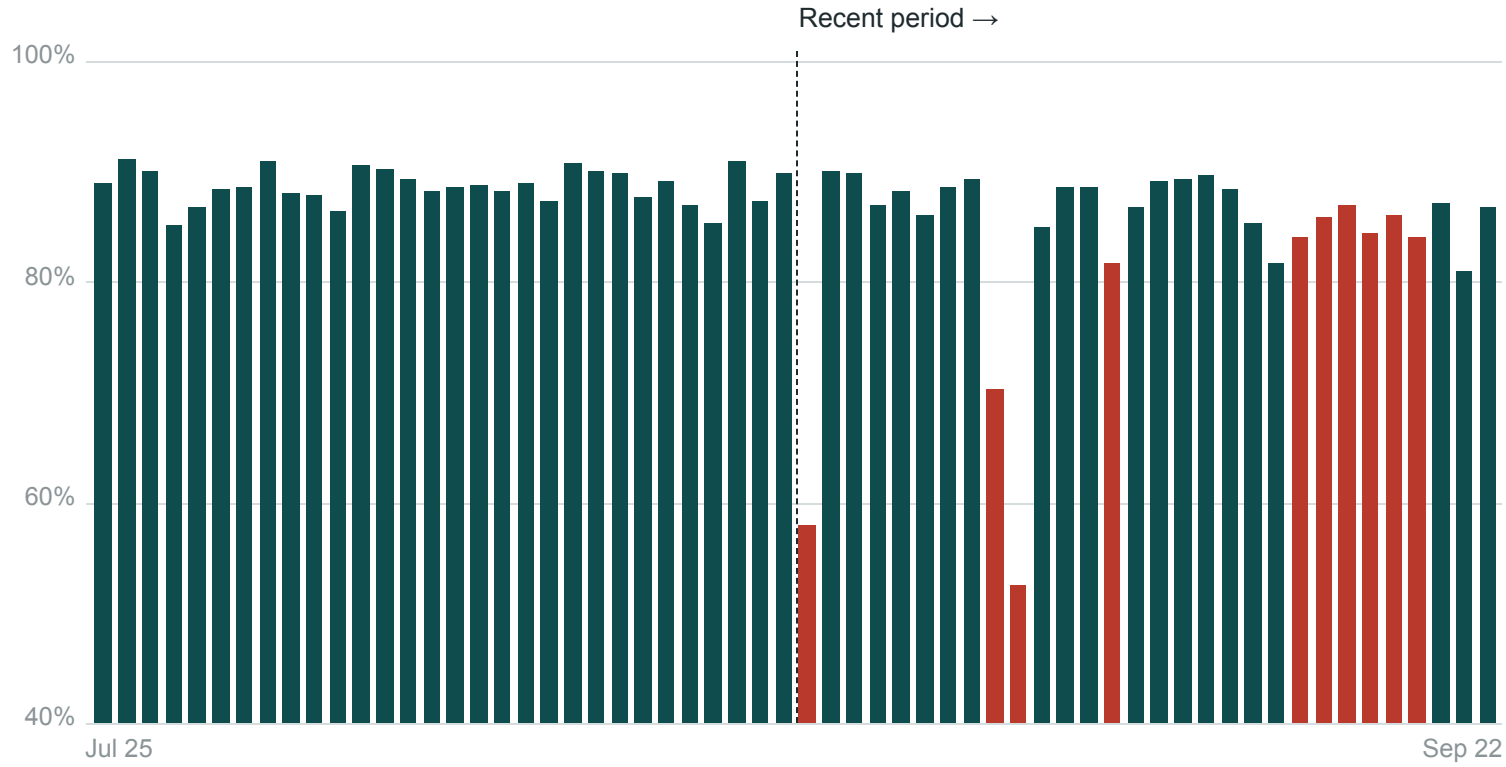
78.8%

Mexican documents (was 84.1%)

89.1%

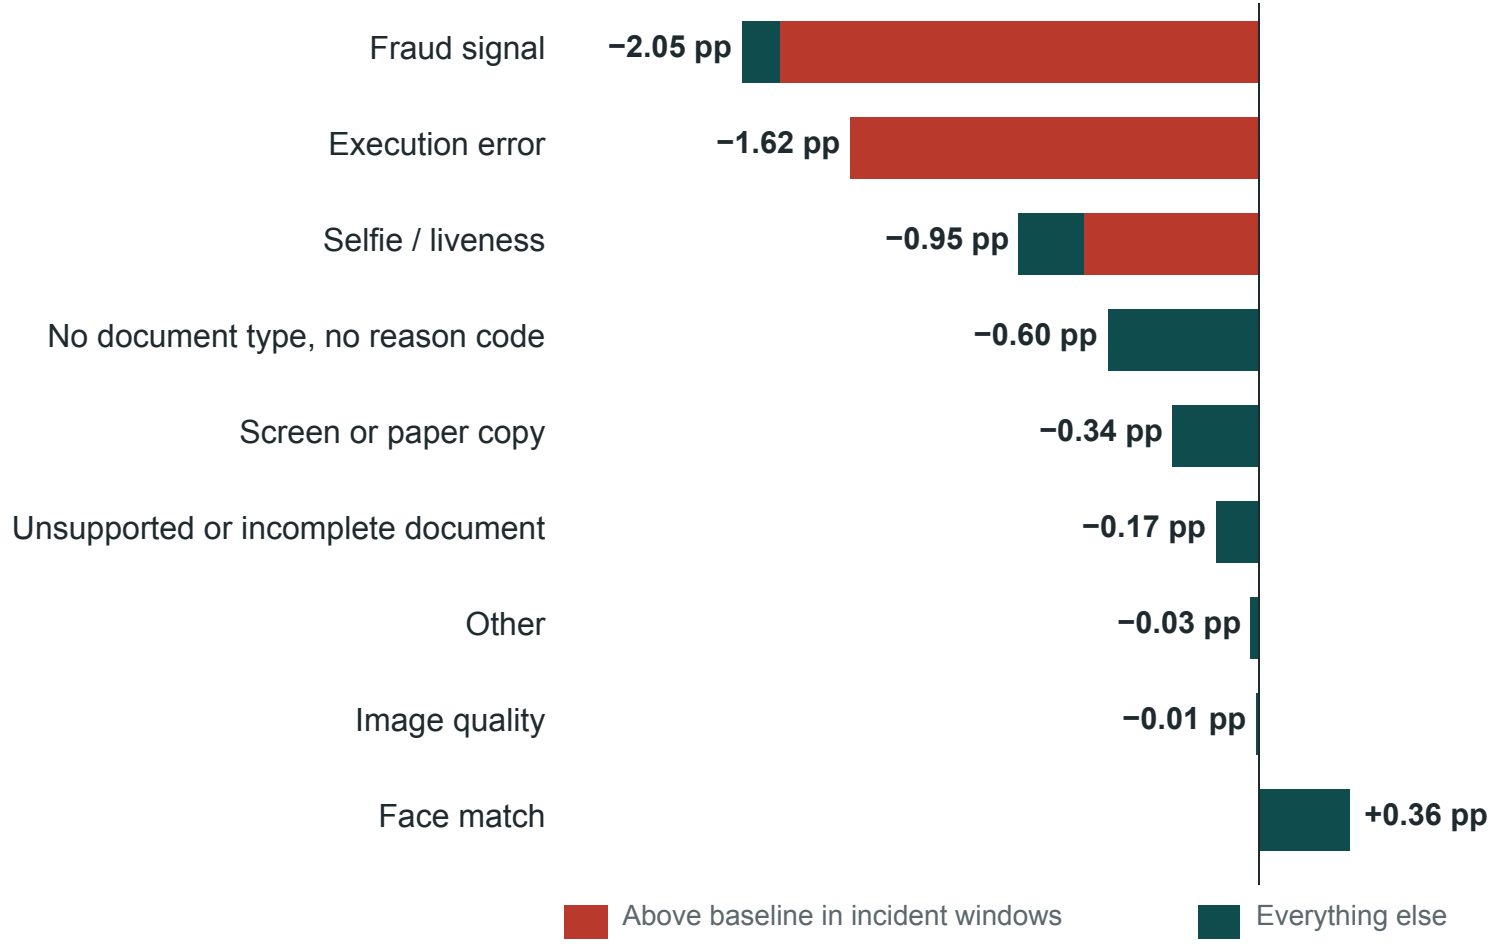
Argentine documents (was 94.9%)

Daily approval rate, all documents (red = incident window days)



Country is the document issuing country. Mix barely moved: Mexican documents were 57.3% of decisions, then 56.8%. Columns start at 40%.

Four reasons explain most of the change; face match improved but remains a large loss



Reading the bars

- Fraud signals: −1.90 pp of −2.05 pp sits in two windows.
- Execution errors: almost all of −1.62 pp is on Sep 1–2.
- Selfie / liveness: −0.69 pp on Sep 1 and 6, −0.26 pp elsewhere.
- No document type, no reason code: −0.60 pp, all outside the windows, so it is a slower, ongoing shift (mostly Mexican documents).
- Face match: 2.77% → 2.41% of decisions. Not a cause of the drop, but a large existing loss.

298 recent unsuccessful decisions had two or more failing checks. A primary reason avoids double counting; it does not mean fixing that reason gets the person approved.

About 1,057 failures sit above baseline in four windows; causes are unconfirmed

Window (2023)	Type	Observed	Expected	Above baseline	Of recent unsuccessful	Approval pts
1. Argentina fraud-signal burst, Aug 24	Suspected additional attack traffic	376	2.9	373	8.9%	−1.48 pp
2. Execution errors at usability, Sep 1–2	Technical failure affecting applicants	407	0.1	407	9.7%	−1.62 pp
3. Liveness quality spike, Sep 1 and Sep 6	Technical failure affecting applicants	174	0.5	173	4.1%	−0.69 pp
4. Mexico INE fraud-signal cluster, Sep 14–19	Fraud-signal cluster without extra volume	111	7.0	104	2.5%	−0.41 pp
Total				**1,057**	**25.2%**	**−4.20 pp**

What is and isn’t known

- Argentina, Aug 24: 660 decisions vs a typical 318; 373 of the flagged passed liveness. Suggests attack traffic; not adjudicated.
- Sep 1–2: usability TECHNICAL_ERROR, liveness passed in all 407, screening passed in 251. Vendor or ARQ integration: unknown.
- Mexico, Sep 14–19: INE flags without extra volume: fraud, a rule change or a new card version.

Vendor status page (read Oct 13, 2026)

Liveness incidents on the vendor’s US data centre: Sep 1 (15:38–16:42 and 18:07–22:44 UTC) and Sep 6 (10:00–17:22 UTC). No entry for Sep 2 or for the fraud-signal dates.

A documented vendor incident is not confirmation of impact on ARQ.

Sensitivity range 1,055–1,130 across window, signature and baseline variants. Removing the 1,068 matching failures gives 86.99%: a descriptive rate, not performance without incidents.

Six hypotheses, each with the step that would settle it

	Hypothesis	Strongest alternative	Next step	Decision it enables
H1	Execution errors on Sep 1–2 became rejections	Vendor fault vs ARQ integration or network	Pull ARQ logs and raw responses; vendor incident report	Fault owner; SLA question; hold-and-retry either way
H2	Liveness failures on Sep 1 and 6 were a vendor fault	ARQ app or SDK release	Release log; vendor list of affected transactions	Escalate; test a retake policy
H3	Aug 24 Argentina burst was coordinated fraud	Vendor false positives during a campaign	Adjudicate a sample; channel and device data	Pre-vendor velocity controls
H4	Sep 14–19 INE flags reflect a vendor rule change	Quieter fraud; new INE card version	Sample review with the vendor; registry check	Review vs auto-decline for INE flags
H5	Mexican "no document type" has several causes	Unsupported documents, capture, configuration, coverage	Review 100 cases with the vendor; config history	Picker vs configuration vs escalation
H6	Face mismatches partly reflect old ID photos	Impostors; selfie lighting	Adjudicate NO_MATCH and NOT_POSSIBLE samples	Review queue or registry face check

Evidence for each (from the data): H1 407 TECHNICAL_ERROR rejections; H2 174 liveness quality failures on two days; H3 376 signals with doubled volume; H4 111 INE signals at normal volume; H5 no document type 11.3% → 14.2% of Mexican decisions; H6 face match 2.41% of recent decisions.

Phase 0, weeks 0–2: measure and stabilise before optimising

Minimum funnel instrumentation

Data & Analytics · Measure · M

Done when: Events for 99% of sessions in a daily reconciliation against vendor transactions; dashboards show first cohort.

Hold and retry on TECHNICAL_ERROR

Engineering · Onboarding platform · Remediate · S

Done when: Zero final rejections with TECHNICAL_ERROR over 4 weeks; 95% of held cases resolved within 24 hours.

Daily, then hourly, alerts per reason code

Data & Analytics · Measure · S

Done when: On replay, alerts fire on the first day of every incident window, with thresholds tuned so prior-period false alerts stay under one a week.

Reconcile the control findings

Compliance · Control · S

Done when: Every record classified (confirmed, explained, data error) with evidence.

Adjudicate a sample of fraud signals

Risk & Fraud · Investigate · S

Done when: At least 50 cases per cluster adjudicated with documented rationale.

Diagnose 'no document type' in Mexico

Product · Onboarding · Investigate · S

Done when: Each sampled case assigned a cause with vendor agreement.

Confirm vendor incident scope

Vendor management · Investigate · S

Done when: Written vendor confirmation or denial per date.

Report the headline with its breakdown

Data & Analytics · Measure · S

Done when: Weekly report adopted by the funnel review.

Effort sizes are assumptions (S ≤ 2 engineer-weeks, M = 2–6). Owners are role-based examples, not named people.

Later phases depend on what Phase 0 finds; platform work is multi-quarter

P1 · Weeks 2–6 · Test the leading fixes

Test one guided retake on unclear selfies

Product · Onboarding
Effort M · confidence low
If: Only if vendor rules permit retakes for these labels.

Test accepted-document guidance in Mexico

Product · Onboarding
Effort M · confidence low
If: Only if the diagnosis points to document choice rather than configuration or vendor coverage.

Velocity and device checks before the vendor call

Risk & Fraud
Effort M · confidence medium
If: Depends on adjudication confirming coordinated fraud.

P2 · Weeks 6–13 · Scale what works

On-device capture quality checks

Engineering · Mobile
Effort M · confidence low
If: Camera-only capture depends on the existing flow and on the cause being uploads.

Pilot manual review of face-match warnings

Operations · KYC review
Effort M · confidence low
If: Requires capacity and quality planning; any threshold change requires risk validation.

Invite recoverable applicants back

Growth & Marketing
Effort S · confidence low
If: Only for appropriate recoverable categories and with consent.

P3 · Next two quarters · Platform work (large)

KYC orchestration and decision layer

Engineering · Onboarding platform
Effort L · confidence medium
If: Depends on what the current architecture already provides.

Government registry checks (INE, RENAPER)

Vendor management
Effort L · confidence low
If: Requires integration and commercial work.

Second verification vendor for failover

Vendor management
Effort L · confidence low
If: Requires integration and commercial work.

Timelines match the dashboard and FAQ: P0 weeks 0–2, P1 weeks 2–6, P2 weeks 6–13, P3 the following two quarters.

Recovery scenarios: +0.9 pp to +4.3 pp on recurring causes, not a forecast

+0.9 pp

Conservative: 233 approvals per 30 days; recent rate would be 84.2%

+2.5 pp

Central: 634 approvals per 30 days; recent rate would be 85.8%

+4.3 pp

Optimistic: 1,093 approvals per 30 days; recent rate would be 87.6%

295

Central: approvals protected per comparable incident (execution errors, liveness spikes)

−580

If all control findings are confirmed: approvals per 30 days that may need reversing (−2.3 pp)

Recovered = observed population × eligibility × reach × return × remediation × subsequent approval

Population (recent period)	Observed	Central factors (E × R × Ret × Fix × Then)	Central	Range
Execution errors (Sep 1–2) (per incident)	407	1.0 × 1.0 × 0.75 × 0.9 × 0.88	242	116–351
Liveness quality failures on incident days (per incident)	174	0.75 × 1.0 × 0.6 × 0.75 × 0.9	53	23–84
Other selfie / liveness failures	486	0.66 × 1.0 × 0.6 × 0.45 × 0.9	78	28–138
Mexico: no document type or unsupported / incomplete document	1,570	0.7 × 1.0 × 1.0 × 0.35 × 0.84	323	107–561
Image quality and screen or paper copies	333	0.85 × 1.0 × 1.0 × 0.45 × 0.85	108	48–162
Face match: WARNING plus a share of NO_MATCH	278	1.0 × 1.0 × 1.0 × 0.45 × 1.0	125	50–232

A target set from these scenarios is a goal, not a prediction. "Eligible" for selfie populations is the observed share with no other failing check; "Then" is the prior-period approval rate for the same countries, except review, where a reviewer approval is final.

What we need, and how we will know it worked

Decisions and resources

- Approve Phase 0 and its owners.
- Privacy review of the event schema (Legal & Privacy).
- Vendor: case access, raw responses, incident reports for Sep 1, 2 and 6.
- Risk and Compliance time: fraud adjudication and control reconciliation.
- Engineering: instrumentation (M), hold-and-retry (S), alerts (S). Sizes to confirm.
- Finance: replace the value model’s market-based presets with ARQ unit economics.

How success is verified

- Zero final rejections with TECHNICAL_ERROR over 4 weeks.
- Alerts replayed on 2023 data fire on the first day of all four windows.
- Each control finding classified as confirmed, explained or data error.
- Experiments by person, outcome = eventual approval within 7 days. Mexico guidance: about 6,314 people per arm for +2 pp (stand-in baseline).
- Weekly: approval rate with fraud-signal, execution-error and warning shares beside it; completion once events exist.

Four things to confirm with the team before acting

	Question	Why it matters	What the answer changes
1	Does each row hold a user's first, latest or final KYC attempt, and does the date mark the submission or the vendor's decision?	The headline is approval among recorded decisions. If rows are latest attempts, earlier failures and retries are invisible and the first-pass rate is lower than shown; if they are final decisions, the warnings may be stale.	The metric definition (v2 to v3 in the registry), the retry hypotheses, and whether the bridge counts people or attempts.
2	Is there any record, even weekly counts, of people who opened KYC and never submitted?	The brief defines the pass rate as people who start and finish. Without starts, completion and abandonment can't be measured, and the 1,587 fewer approvals from lower volume could be fewer starts, more abandonment, or both.	Whether Phase 0 instrumentation is a prerequisite or a check on existing data; the reading of the volume effect; the completion metric's status.
3	What happens today after a TECHNICAL_ERROR or a liveness failure: does the user see the vendor's outcome, get a retake, or enter a review queue?	Hold-and-retry and the guided retake assume no retry exists. In the central scenario they recover 242 approvals per comparable execution-error incident and 78 per 30 days on steady liveness failures.	If a retry already exists, those recoverable counts shrink and the question becomes why retries fail; the Phase 0 hold-and-retry becomes a configuration check.
4	Which vendor product tier and data centre does ARQ use, and did ARQ receive incident notices for Sep 1, 2 and 6, 2023?	The vendor's status history shows liveness incidents on Sep 1 and 6 on its US data centre and nothing for Sep 2. The liveness spike is 173 failures above baseline and the execution errors 407.	On another data centre the liveness attribution loses its vendor evidence; with notices, the execution errors have an owner and an SLA claim.

The analysis stands on the extract alone. These answers decide how far each conclusion carries, not whether the work was done.

Appendix: evidence and method

A1 Definitions · A2 Validation · A3 Composition · A4 Mexico · A5 Incident sensitivity · A6 Experiments · A7–A8 Value model · A9 Operating model · A10 Control findings · A11 Vendor research

A1 · Metric definitions: what is available and what needs new events

Metric	Status	Definition and window	Gap or caveat
KYC starts	Unavailable	People who open the KYC flow. Window: daily, by country.	kyc_started with person id, session id, timestamp, country, app version, channel
Recorded decisions	Available	User records with a final decision in the extract (one per user_reference). Window: per period.	Whether a record is a first, latest or final attempt is not stated
Approval rate among recorded decisions	Available	Approved (PASSED or APPROVED) ÷ recorded decisions. The headline metric of this analysis. Window: per period, by decision date.	Not the start-to-finish completion rate: starts, abandonment and attempts are missing
Rejection rate	Available	REJECTED ÷ recorded decisions. Window: per period.	—
Unresolved warning rate	Available	WARNING ÷ recorded decisions. Later review outcome not in the extract. Window: per period.	Review dispositions needed
Completion rate (start to terminal state)	Unavailable	Starts that reach a terminal state (approved, finally rejected, or abandoned for 7 days) within 7 days of start. Window: 7 days after start.	kyc_started, step events, terminal decision, abandonment rule
Eventual approval within 7 days	Unavailable	Starts approved on any attempt within 7 days of the first start. Window: 7 days, cohort by start date.	Person id across attempts, attempt number, timestamps, review dispositions
Abandonment rate	Unavailable	Starts with no submission within 7 days. Window: 7 days.	kyc_started and kyc_submitted events
Retry rate and retry approval	Unavailable	Share of unsuccessful first attempts that retry, and their approval. Window: 7 days.	Attempt number per person
Manual-review outcomes	Unavailable	Approval and turnaround of cases sent to review. Window: per week.	Review queue events and dispositions
Time to completion	Unavailable	Median minutes from start to terminal decision. Window: per week.	Timestamps with time of day on start and decision
Execution-error rate	Available	Unsuccessful decisions whose primary reason is TECHNICAL_ERROR ÷ recorded decisions. Window: per period.	—
Fraud-signal rate	Available	Unsuccessful decisions whose primary reason is a manipulation or fake label ÷ recorded decisions. Not adjudicated. Window: per period.	Adjudication of a sample

A2 · Data validation and outcome treatment

Check	Result	Status	Treatment
Records in each extract	52,075 summary / 52,075 details	Pass	Equal row counts.
Unique user_reference	52,075 / 52,075	Pass	No duplicate users in either file.
Join coverage on user_reference	52,075 matched, 0 + 0 unmatched	Pass	Complete one-to-one join.
Date coverage	2023-07-25 to 2023-09-22: 60 of 60 calendar days	Pass	No missing days.
Timestamp precision	52,075 of 52,075 at midnight UTC	Limitation	Dates only. Intraday outage windows cannot be matched to records, and the extract does not say which event the date marks (submission, decision or export).
Final outcome categories (summary.decision_type)	PASSED 44,825, REJECTED 6,772, WARNING 474, APPROVED 4	Pass	PASSED and APPROVED = approved; REJECTED = rejected; WARNING = unresolved (not approved, not a confirmed rejection).
decision_type vs details.decision_label	PASSED vs REJECTED: 1; PASSED vs OK: 8	Note	summary.decision_type is treated as the outcome. The 8 'OK' labels are label-only; 1 record shows PASSED in the summary but REJECTED in details and is counted as approved, pending reconciliation.
Non-standard capability value 'PASSES'	3 values	Note	Treated as PASSED.
Blank watchlist_screening_decision	6,606 blank, 1 NOT_EXECUTED	Note	Shown as 'not recorded'. Mostly on rejected records; a blank field does not prove screening never ran elsewhere.
Blank data_type (no document type returned)	3,781 records	Note	A symptom with several possible causes (unsupported document, capture, configuration, classification coverage).
Birth year extracted	48,286 present, 3,789 blank; range 1909–2022	Limitation	3,769 of the 3,789 blanks are rejected records, so ages describe applicants whose document was read, not all applicants. Extreme years (e.g. 2022) suggest extraction errors.
Country field	MEX / ARG only (data_issuing_country)	Note	This is the document's issuing country, not necessarily residence or account market.
Attempts per user	1 record per user_reference	Limitation	The extract does not say whether a record is the first, latest or final attempt; retries and later outcomes are not visible.
Raw files unchanged	summary 67e4190909b7…, details e0762f377a41…	Pass	SHA-256 prefixes, recomputed on every build.

A3 · What the not-approved share is made of

Group	Class	Prior	Recent	Of decisions, prior	Of decisions, recent	Of unsuccessful, recent
Execution error	Execution error	1	407	0.00%	1.62%	9.7%
Fraud signal	Fraud signal (not adjudicated)	88	598	0.33%	2.38%	14.2%
No document type, no reason code	Unexplained or identity mismatch	945	1,034	3.51%	4.11%	24.6%
Unsupported or incomplete document	Capture or document issue	528	536	1.96%	2.13%	12.8%
Image quality	Capture or document issue	186	175	0.69%	0.70%	4.2%
Screen or paper copy	Capture or document issue	78	158	0.29%	0.63%	3.8%
Selfie / liveness	Capture or document issue	451	660	1.68%	2.63%	15.7%
Face match	Unexplained or identity mismatch	747	606	2.77%	2.41%	14.4%
Other	Unexplained or identity mismatch	21	27	0.08%	0.11%	0.6%
Fraud signal (not adjudicated)	Class total	88	598	0.33%	2.38%	14.2%
Execution error	Class total	1	407	0.00%	1.62%	9.7%
Capture or document issue	Class total	1,243	1,529	4.61%	6.08%	36.4%
Unexplained or identity mismatch	Class total	1,713	1,667	6.36%	6.63%	39.7%

Not-approved share: 11.3% prior, 16.7% recent. 281 recent decisions are unresolved warnings, not confirmed rejections.

A4 · Mexico: "no document type" is a symptom with several outcomes

Outcome for Mexican decisions with no document type	Prior	Recent	Share, prior	Share, recent	Change
No reason code	945	1,034	6.13%	7.24%	+1.11 pp
Unsupported document type	245	264	1.59%	1.85%	+0.26 pp
Missing mandatory data	203	204	1.32%	1.43%	+0.11 pp
Missing page or part	80	68	0.52%	0.48%	−0.04 pp
Image quality	186	175	1.21%	1.23%	+0.02 pp
Screen or paper copy	59	116	0.38%	0.81%	+0.43 pp
Execution error	1	96	0.01%	0.67%	+0.67 pp
Selfie, face or other	28	58	0.18%	0.41%	+0.22 pp
Approved	0	19	0.00%	0.13%	+0.13 pp

Document type returned (Mexico)	Approval prior	Approval recent
Voter ID (INE)	97.7%	93.9%
National ID card	91.0%	87.0%
No document type returned	0.0%	0.9%
Passport	97.8%	95.6%
Residence permit	97.6%	94.6%
Driving licence	97.0%	92.5%
Other	100.0%	93.8%

Argentine documents: 0 usability rejections or warnings in the extract; Mexican: 1,442. Possible causes include documents outside the accepted list, capture quality, issuing-country mismatch, configuration and vendor coverage.

A5 · Incident attribution: sensitivity and vendor evidence

Window	Base	Window ±1 day	Broader signature	Recent baseline	Old method (no baseline)
1. Argentina fraud-signal burst	373	378	372	371	376
2. Execution errors at usability	407	407	454	407	407
3. Liveness quality spike	173	174	171	174	174
4. Mexico INE fraud-signal cluster	104	102	133	103	111
Total	1,057				1,068

How to read it

- Base: failures matching the window's signature, minus the count expected at the prior-period rate for the window's volume.
- Range across the four baseline-based variants: 1,055–1,130.
- The old method attributes every matching failure on the day, including ones that would have happened anyway.

Date	UTC	Vendor status page entry
2023-09-01	15:38–16:42	Liveness returned 'Unable to see your face properly' for all ID verification transactions.
2023-09-01	18:07–22:44	Same liveness message. Liveness restored 20:59; Authentication restored for new users only; resolved 22:44.
2023-09-06	10:00–17:22	Instability; workaround 11:25; liveness restored 16:21; resolved 17:22.
2023-09-13	13:23–13:44	Retry surge after a fix: V2 scans unavailable; Jumio Go customers may have seen extra 'unsupported document' rejections; Hybrid delays.
2023-09-01 / 09-02	—	No status entry mentions usability, TECHNICAL_ERROR or skipped document checks, and none is listed for Sep 2.
2023-08-24, 09-14 to 09-19	—	No incidents listed for August 2023; only scheduled maintenance in Singapore during Sep 14–19.

A6 · Experiment and validation designs

	Mexico document guidance	Guided liveness retake
Kind	Optimisation experiment	Optimisation experiment
Hypothesis	Showing Mexican applicants which documents are accepted, with pictures, before capture raises eventual approval.	Offering one guided retake when liveness returns BAD_QUALITY or FACE_NOT_FULLY_VISIBLE raises eventual approval without raising fraud.
Population	People starting KYC with a Mexican document. Exclude staff and test accounts, and people already approved.	People whose first liveness result is BAD_QUALITY or FACE_NOT_FULLY_VISIBLE and who have no fraud signal. Exclude ID_USED_AS_SELFIE and any manipulation label.
Assignment unit	Person (person_id), assigned at first kyc_started, sticky across attempts.	Person, assigned at the first eligible liveness failure, sticky.
Primary outcome	Eventual approval within 7 days of first start.	Eventual approval within 7 days of the failure.
Window	7 days per person plus 7 days to mature the last cohort.	7 days plus 7 days to mature.
Minimum effect	+2.0 pp on eventual approval.	+10 pp on eventual approval in the eligible group.
Sample	Using the recent submission-based Mexican approval rate (78.8%) as a stand-in baseline, about 6,314 people per arm (95% confidence, 80% power). At about 476 Mexican decisions a day that is roughly 27 days at a 50/50 split. Recompute once start-based data exists.	Assuming about 5% of the eligible group is approved later without a retake (not measurable yet), about 141 people per arm. At about 16 eligible people a day (recent non-incident rate) that is roughly 18 days. Needs the real baseline from instrumentation.
Guardrails	Start-to-submit rate, time to decision, fraud-signal rate, support contacts.	Fraud-signal rate and confirmed fraud among retake approvals, liveness spoof labels, vendor cost per approval, time to decision.
Rules	Start at 10% for 3 days to check events, then 50/50. Expand if the primary outcome improves by at least the minimum and guardrails hold. Roll back if start-to-submit falls more than 2 pp or fraud signals rise more than 0.5 pp.	Needs Risk sign-off before launch. Start 50/50. Roll back on any confirmed spoof or a fraud-signal rate among retake approvals above control.

A7 · Value model inputs: market-based presets, not ARQ figures

Input	Low	Mid	High	Kind	Basis	Sources
Value per activated customer (USD)	$13	$21	$53	Derived	ARQ press figures; Nu México revenue and margin	S4, S5, S8, S9
Activation rate after approval (%)	38.2%	56.3%	83.0%	Proxy	Liftoff LATAM and fintech activation; Nu activity rate	S6, S7
Share of applicants from paid channels (%)	29.0%	40.0%	50.0%	Proxy + assumption	Nu organic and word-of-mouth shares	S6
Acquisition cost per paid applicant (USD)	$8.74	$17.96	$29.93	Derived	Liftoff cost per install and per registration	S7
Billable vendor transactions per decision (count)	1	1.2	1.5	Assumption	Assumption; replace with invoiced transactions ÷ decisions	S2
Vendor cost per billable transaction (USD)	$1.44	$1.85	$3	List price / estimate	Veriff and Sumsub list prices; Vendr estimate for Jumio	S1, S2, S3
Cost per manual review (USD)	$0.83	$5.72	$10.79	Derived	Mexico analyst salaries × handling time	S10, S11, S13, S14
Implementation cost of the recurring-cause actions (USD)	$24,400	$31,600	$85,000	Derived	Roadmap effort × Mexico engineering cost	S12, S13, S14
Run cost per month (USD)	$510	$2,010	$3,860	Derived	Review cases and retake vendor calls, central scenario	S1, S2, S3, S10, S11

Sources: S1 Veriff · S2 Sumsub · S3 Vendr · S4 Nu Holdings · S5 Nu Holdings · S6 Nu Holdings · S7 Liftoff, via MarTech Series · S8 Bloomberg Línea · S9 Passageiro de Primeira · S10 TalentUp · S11 Paylab · S12 Levels.fyi · S13 Revelo · S14 Banco de México. Full figures, periods and links are in the dashboard (Plan & value).

A8 · Value model results by preset, central scenario

Result (per 30 days unless stated)	Low	Mid	High	Cautious
Activated customers from recoverable approvals, per 30 days	242	357	526	242
First-year value of one 30-day cohort	$3,148	$7,496	$27,890	$3,148
Same, if every 30 days repeats for a year	$38,306	$91,199	$339,324	$38,306
Break-even first-year value per activated customer	$11	$13	$21	$45
Review cost per 30 days for the review-pilot population	$231	$1,590	$3,000	$3,000
Historical vendor spend on recent not-approved decisions	$6,049	$9,326	$18,905	$6,049
Vendor spend avoidable on suspected extra attack traffic	$492	$759	$1,539	$492
Historical paid acquisition spend on recent not-approved applicants	$10,648	$30,180	$62,868	$10,648
Pays back within the first year? (value per activated customer vs break-even)	Yes: $13 vs $11	Yes: $21 vs $13	Yes: $53 vs $21	No: $13 vs $45

The first-year value of a cohort is earned over the following year, not recognised in the next 12 months. Past vendor and paid-acquisition spend on not-approved applicants is spent; it is not a saving unless the cause is removed.

A9 · Operating model: owned metrics with versions, lineage and adoption

Metric	Status	Owner (role)	Version	Lineage
Approval rate among recorded decisions	Available	Product · Onboarding	2.0	KYC_summary. decision_type
Completion rate (start to terminal state)	Unavailable	Product · Onboarding	1.0	Needs new events
Eventual approval within 7 days	Unavailable	Product · Onboarding	1.0	Needs new events
Manual-review outcomes	Unavailable	Operations · KYC review	1.0	Needs new events
Execution-error rate	Available	Engineering · Onboarding platform	1.0	KYC_summary. decision_type
Fraud-signal rate	Available	Risk & Fraud	1.0	KYC_summary. decision_type

Ownership rules

- One accountable owner per metric, effective-dated.
- Definition changes need owner and Data approval; old versions stay queryable.
- Past results stay attributed to the owner at the time.

Adoption success measures

Share of KYC actions with an owner and a current status; share of metrics with an approved definition; time from alert to acknowledged; decisions taken in the review that cite the platform. Page views are not a success measure.

This analysis is version 2.0 of the headline metric: v1.0 called it "pass rate"; v2.0 renamed it and defined completion separately.

A10 · Control findings in the extract: to reconcile, not confirmed failures

Finding	Records	Other explanations	How to validate
Approved records with birth year 2006 or later	1,223	Extraction error; a policy that allows minors in some product; an approval record that never became an enabled account.	Compare with raw vendor responses and document images for a sample; check the age policy per country; check account state.
Approved records with no screening result in the extract	86	Screening ran in another system or a later step; field not exported.	Look up screening records for these users in the screening system; confirm whether an account was enabled.
Approved records with a watchlist WARNING	69	Warnings cleared in manual review that the extract doesn't show.	Retrieve review dispositions; confirm who cleared each case and when.
Outcome disagrees between the two files	9	Export mapping difference.	Reconcile against the system of record.
Non-standard capability values	3	Manual edits or an export bug.	Fix at source; add schema checks to the extract.

If an extracted birth year of 2006 or later is accurate, the person was under 18 during 2023. The applicable eligibility policy and the data quality both still need confirming.

A11 · Vendor research (Oct 13, 2026): availability is not effectiveness

Capability	2023 problem	Documented	Prerequisites	Not verified for ARQ
Classification	No document type returned	A separate capability that returns the document type, including UNSUPPORTED.	None stated.	Whether ARQ's workflow already returns it.
Usability labels	Missing mandatory data; automation level	UNSUPPORTED_DOCUMENT_TYPE is a WARNING; MISSING_MANDATORY_DATAPOINTS 'will only happen for Jumio Go transactions'; TECHNICAL_ERROR is NOT_EXECUTED.	None stated.	Whether ARQ had manual review elsewhere in its own process.
Liveness labels and tiers	Liveness quality spike	BAD_QUALITY is a WARNING; FACE_NOT_FULLY_VISIBLE is REJECTED; Standard or Premium type.	Premium 'requires additional configuration for your tenant'; technology depends on the product purchased.	Which tier and labels ARQ receives.
Cross Transaction Risk	Fraud-signal clusters	Matches ID data, faces and image parts such as backgrounds across transactions; informational by default.	Enabled by contacting Jumio Support. Press release of Apr 22, 2025 announced general availability.	Commercial terms; effect on ARQ's traffic.
Result retrieval and retries	Execution errors	After a callback, retry retrieval from 40 seconds with growing intervals, at most 10 consecutive failures.	None.	How ARQ handles TECHNICAL_ERROR today.
Mobile SDK capture features	No document type; capture quality	ML document and country classifier (4.9.0); unsupported-document check (4.11.0); virtual-camera injection detection (4.14.0); controls to limit manual capture (4.18.0).	SDK upgrade; configuration details for 4.18.0 controls not documented; gallery upload not documented.	ARQ's current SDK version and capture flow.
Mexico INE registry checks	INE fraud signals; face mismatch	INE Lista Nominal validates card fields. INE Data & Face Validation compares with INE records.	Data & Face needs a direct contract with INE, onboarding with FIMPE, dedicated middleware, about 4–5 days setup plus 2 days certification per client.	Cost, coverage and effect.
Argentina RENAPER data and biometric	Fraud signals; face mismatch	Data check and biometric (selfie vs registry) check.	PII fields need an updated data-processing agreement and are disabled by default; enabled by Jumio.	Cost and effect.